

## MODULE 2 : INPUT/OUTPUT ORGANIZATION

### ACCESSING I/O DEVICES

- There are 2 ways to deal with I/O devices (Figure 4.1).

#### 1) Memory mapped I/O

- Memory and I/O devices share a common address-space.
- Any data-transfer instruction (like Move, Load) can be used to exchange information.
- For example, Move DATAIN, R0; this instruction reads data from DATAIN (input-buffer associated with keyboard) & stores them into processor-register R0.

#### 2) In *I/O mapped I/O*, memory and I/O address-spaces are different.

- Special instructions named IN and OUT are used for data transfer.
- Advantage of separate I/O space: I/O devices deal with fewer address-lines.




Figure 4.1 A single-bus structure.

### I/O Interface for an Input Device

- Address decoder: decodes address sent on bus, so as to enable input-device (Figure 4.2).
- Data register: holds data being transferred to or from the processor.
- Status register: contains information relevant to operation of I/O device.
- Address decoder, data- and status-registers, and control-circuitry required to coordinate I/O transfers constitute device's interface-circuit.

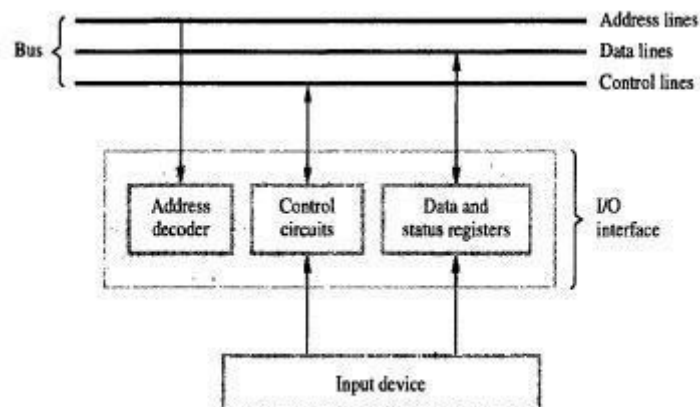


Figure 4.2 I/O interface for an input device.

### MECHANISMS USED FOR INTERFACING I/O DEVICES

#### 1) Program Controlled I/O

- Processor repeatedly checks a status-flag to achieve required synchronization between processor & input/output device. (We say that the processor polls the device).
- Main drawback: The processor wastes its time in checking the status of the device before actual data transfer takes place.

#### 2) Interrupt I/O

- Synchronization is achieved by having I/O device send a special signal over bus whenever it is ready for a data transfer operation.

#### 3) Direct Memory Access (DMA)

- This involves having the device-interface transfer data directly to or from the memory without continuous involvement by the processor.

### INTERRUPTS

- I/O device initiates the action instead of the processor. This is done by sending a special hardware signal to the processor called as **interrupt**(INTR), on the interrupt-request line.
- The processor can be performing its own task without the need to continuously check the I/O device.
- When device gets ready, it will "alert" the processor by sending an interrupt-signal (Figure 4.5).
- The routine executed in response to an interrupt-request is called ISR(Interrupt Service Routine).
- Once the interrupt-request signal comes from the device, the processor has to inform the device that its request has been recognized and will be serviced soon. This is indicated by a special control signal on the bus called **interrupt-acknowledge**(INTA).

#### Difference between subroutine & ISR

- A subroutine performs a function required by the program from which it is called.  
However, the ISR may not have anything in common with the program being executed at the time the interrupt-request is received. Before starting execution of ISR, any information that may be altered during the execution of that routine must be saved. This information must be restored before the interrupted-program resumed.
- Another difference is that an interrupt is a mechanism for coordinating I/O transfers whereas a subroutine is just a linkage of 2 or more function related to each other.

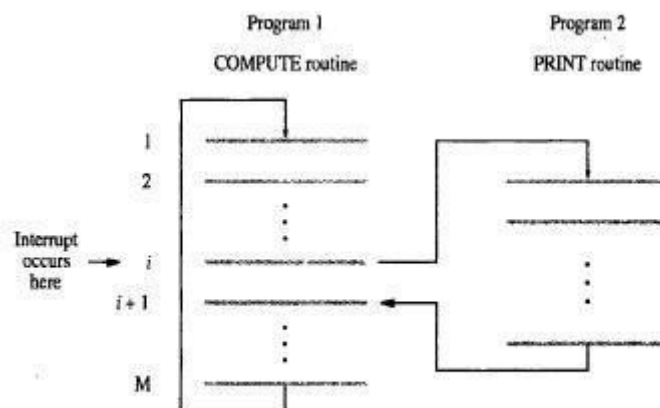


Figure 4.5 Transfer of control through the use of interrupts.

- The speed of operation of the processor and I/O devices differ greatly. Also, since I/O devices are manually operated in many cases (like pressing a key on keyboard), there may not be synchronization between the CPU operations and I/O operations with reference to CPU clock. To cater to the different needs of I/O operations, 3 mechanisms have been developed for interfacing I/O devices. 1) Program controlled I/O 2) Interrupt I/O 3) Direct memory access (DMA).

## Computer Organization

- Saving registers increases the delay between the time an interrupt request is received and the start of execution of the ISR. This delay is called interrupt latency.
- Since interrupts can arrive at any time, they may alter the sequence of events. Hence, facility must be provided to enable and disable interrupts as desired.
- Consider the case of a single interrupt request from one device. The device keeps the interrupt request signal activated until it is informed that the processor has accepted its request. This activated signal, if not deactivated may lead to successive interruptions, causing the system to enter into an infinite loop.

### INTERRUPT HARDWARE

- An I/O device requests an interrupt by activating a bus-line called interrupt-request(IR).
- A single IR line can be used to serve „n“ devices (Figure 4.6).
- All devices are connected to IR line via switches to ground.
- To request an interrupt, a device closes its associated switch. Thus, if all IR signals are inactive(i.e. if all switches are open), the voltage on the IR line will be equal to  $V_{dd}$ .
- When a device requests an interrupt by closing its switch, the voltage on the line drops to 0, causing the INTR received by the processor to goto 1.
- The value of INTR is the logical OR of the requests from individual devices
- A special gate known as open-collector or open-drain are used to drive the INTR line.
- Resistor R is called a *pull-up resistor* because it pulls the line voltage up to the high-voltage state when the switches are open.




Figure 4.6 An equivalent circuit for an open-drain bus used to implement a common interrupt-request line.

### ENABLING & DISABLING INTERRUPTS

- To prevent the system from entering into an infinite-loop because of interrupt, there are 3 possibilities:
  - 1) The first possibility is to have the processor-hardware ignore the interrupt-request line until the execution of the first instruction of the ISR has been completed.
  - 2) The second option is to have the processor automatically disable interrupts before starting the execution of the ISR.
  - 3) In the third option, the processor has a special interrupt-request line for which the interrupt-handling circuit responds only to the leading edge of the signal. Such a line is said to be edge-triggered.
- Sequence of events involved in handling an interrupt-request from a single device is as follows:
  - 1) The device raises an interrupt-request.
  - 2) The program currently being executed is interrupted.
  - 3) All interrupts are disabled(by changing the control bits in the PS).
  - 4) The device is informed that its request has been recognized, and in response, the device deactivates the interrupt-request signal.
  - 5) The action requested by the interrupt is performed by the ISR.
  - 6) Interrupts are enabled again and execution of the interrupted program is resumed.

## HANDLING MULTIPLE DEVICES

### Polling

- Information needed to determine whether a device is requesting an interrupt is available in its status-register.
- When a device raises an interrupt-request, it sets IRQ bit to 1 in its status-register (Figure 4.3).
- KIRQ and DIRQ are the interrupt-request bits for keyboard & display.
- Simplest way to identify interrupting device is to have ISR poll all I/O devices connected to bus.
- The first device encountered with its IRQ bit set is the device that should be serviced. After servicing this device, next requests may be serviced.
- Main advantage: Simple & easy to implement.

Main disadvantage: More time spent polling IRQ bits of all devices (that may not be requesting any service).




Figure 4.3 Registers in keyboard and display interfaces.

	Move	#LINE,R0	Initialize memory pointer.
WAITK	TestBit	#0,STATUS	Test SIN.
	Branch=0	WAITK	Wait for character to be entered.
	Move	DATAIN,R1	Read character.
WAITD	TestBit	#1,STATUS	Test SOUT.
	Branch=0	WAITD	Wait for display to become ready.
	Move	R1,DATAOUT	Send character to display.
	Move	R1,(R0)+	Store character and advance pointer.
	Compare	#\$0D,R1	Check if Carriage Return.
	Branch≠0	WAITK	If not, get another character.
	Move	#\$0A,DATAOUT	Otherwise, send Line Feed.
	Call	PROCESS	Call a subroutine to process the input line.

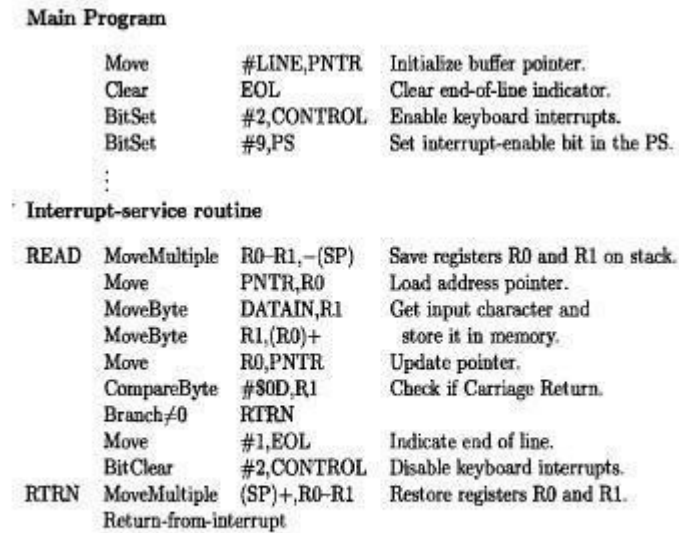
Figure 4.4 A program that reads one line from the keyboard, stores it in memory buffer, and echoes it back to the display.

### Vectored Interrupts

- A device requesting an interrupt identifies itself by sending a special-code to processor over bus. (This enables processor to identify individual devices even if they share a single interrupt-request line).
- The code represents starting-address of ISR for that device.
- ISR for a given device must always start at same location.
- The address stored at the location pointed to by interrupting-device is called the interrupt-vector.
- Processor
  - loads interrupt-vector into PC &
  - executes appropriate ISR
- Interrupting-device must wait to put data on bus only when processor is ready to receive it.
- When processor is ready to receive interrupt-vector code, it activates INTA line.
- I/O device responds by sending its interrupt-vector code & turning off the INTR signal.

### CONTROLLING DEVICE REQUESTS


- There are 2 independent mechanisms for controlling interrupt requests.
- At device-end, an interrupt-enable bit in a control register determines whether device is allowed to generate an interrupt request.
- At processor-end, either an interrupt-enable bit in the PS register or a priority structure determines whether a given interrupt-request will be accepted.



**Figure 4.9** Using interrupts to read a line of characters from a keyboard via the registers in Figure 4.3.

### INTERRUPT NESTING

- A multiple-priority scheme is implemented by using separate INTR & INTA lines for each device
- Each of the INTR lines is assigned a different priority-level (Figure 4.7).
- Priority-level of processor is the priority of program that is currently being executed.
- During execution of an ISR, interrupt-requests will be accepted from some devices but not from others depending upon device's priority.
- Processor accepts interrupts only from devices that have priority higher than its own.
- At the time of execution of an ISR for some device is started, priority of processor is raised to that of the device
- Processor's priority is encoded in a few bits of processor-status (PS) word. This can be changed by program instructions that write into PS. These are called *privileged instructions*.
- Privileged-instructions can be executed only while processor is running in supervisor-mode.
- Processor is in supervisor-mode only when executing operating-system routines. (An attempt to execute a privileged-instruction while in the user-mode leads to a special type of interrupt called a *privileged exception*).



**Figure 4.7** Implementation of interrupt priority using individual interrupt-request and acknowledge lines.

### SIMULTANEOUS REQUESTS

- INTR line is common to all devices (Figure 4.8).
- INTA line is connected in a daisy-chain fashion such that INTA signal propagates serially through devices.
- When several devices raise an interrupt-request and INTR line is activated, processor responds by setting INTA line to 1. This signal is received by device 1.
- Device 1 passes signal on to device 2 only if it does not require any service.
- If device 1 has a pending-request for interrupt, it blocks INTA signal and proceeds to put its identifying code on data lines.
- Device that is electrically closest to processor has highest priority.
- Main advantage: This allows the processor to accept interrupt-requests from some devices but not from others depending upon their priorities.

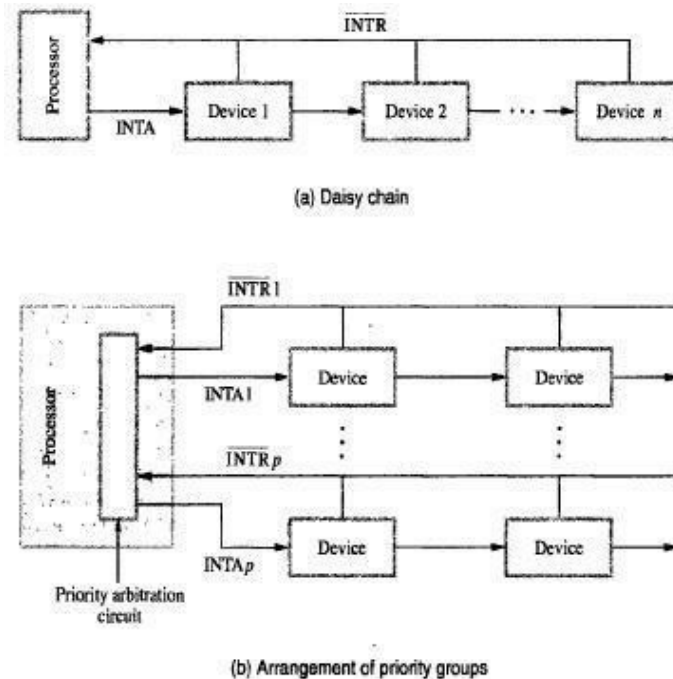


Figure 4.8 Interrupt priority schemes.

### EXCEPTIONS

- An interrupt is an event that causes
  - execution of one program to be suspended &
  - execution of another program to begin.
- *Exception* refers to any event that causes an interruption. I/O interrupts are one example of an exception.
- Computers use a variety of techniques to ensure that all hardware-components are operating properly. For e.g. many computers include an error-checking code in main-memory which allows detection of errors in stored-data.
- If an error occurs, control-hardware detects it & informs processor by raising an interrupt.
- When exception processing is initiated (as a result of errors), processor
  - suspends program being executed &
  - starts an ESR(Exception Service Routine). This routine takes appropriate action to recover from the error to inform user about it.



## Computer Organization

- Debugger
  - helps programmer find errors in a program and
  - uses exceptions to provide 2 important facilities: 1) Trace & 2) Breakpoints
- When a processor is operating in trace-mode, an exception occurs after execution of every instruction (using debugging-program as ESR).
- Debugging-program enables user to examine contents of registers (AX, BX), memory-locations and so on.
- On return from debugging-program,
  - next instruction in program being debugged is executed,
  - then debugging-program is activated again.
- Breakpoints provide a similar facility except that program being debugged is interrupted only at specific points selected by user. An instruction called Trap(or Software interrupt) is usually provided for this purpose.

### Privilege Exception

- To protect OS of computer from being corrupted by user-programs, certain instructions can be executed only while processor is in supervisor-mode. These are called *privileged instructions*.
- For e.g. when the processor is running in user-mode, it will not execute an instruction that changes priority-level of processor.
- An attempt to execute such an instruction will produce a privilege-exception. As a result, processor switches to supervisor-mode & begins to execute an appropriate routine in OS.

## DIRECT MEMORY ACCESS (DMA)

- The transfer of a block of data directly between an external device & main memory without continuous involvement by processor is called as *DMA*.
- DMA transfers are performed by a control-circuit that is part of I/O device interface. This circuit is called as a *DMA controller* (Figure 4.19).
- DMA controller performs the functions that would normally be carried out by processor
- In controller, 3 registers are accessed by processor to initiate transfer operations (Figure 4.18):
  - 1) Two registers are used for storing starting-address & word-count
  - 2) Third register contains status- & control-flags
- The R/W bit determines direction of transfer.
  - When R/W=1, controller performs a read operation(i.e. it transfers data from memory to I/O), Otherwise it performs a write operation (i.e. it transfers data from I/O device to memory).
- When Done=1, controller
  - completes transferring a block of data &
  - is ready to receive another command.
- When IE=1, controller raises an interrupt after it has completed transferring a block of data (IE=Interrupt Enable).
- Finally, when IRQ=1, controller requests an interrupt. (Requests by DMA devices for using the bus are always given higher priority than processor requests).
- There are 2 ways in which the DMA operation can be carried out:
  - 2) In one method, processor originates most memory-access cycles. DMA controller is said to "steal" memory cycles from processor. Hence, this technique is usually called *cycle stealing* or *data without any interruption* (this is known as *block mode* for bus mode).




Figure 4.18 Registers in a DMA interface.

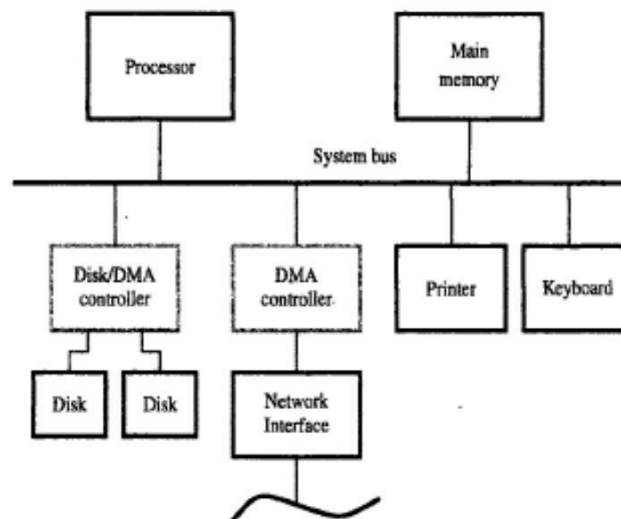


Figure 4.19 Use of DMA controllers in a computer system.



## Computer Organization

### BUS ARBITRATION

- The device that is allowed to initiate data transfers on bus at any given time is called *bus-master*.
- There can be only one bus master at any given time.
- *Bus arbitration* is the process by which next device to become the bus-master is selected and bus-mastership is transferred to it.
- There are 2 approaches to bus arbitration:
  - 1) In centralized arbitration, a single bus-arbiter performs the required arbitration.
  - 2) In distributed arbitration, all device participate in selection of next bus-master.

### CENTRALIZED ARBITRATION

- A single bus-arbiter performs the required arbitration (Figure: 4.20 & 4.21).
- Normally, processor is the bus master unless it grants bus mastership to one of the DMA controllers.
- A DMA controller indicates that it needs to become busmaster by activating Bus-Request line(BR).
- The signal on the BR line is the logical OR of bus-requests from all devices connected to it.
- When BR is activated, processor activates Bus-Grant signal(BG1) indicating to DMA controllers that they may use bus when it becomes free. (This signal is connected to all DMA controllers using a daisy-chain arrangement).
- If DMA controller-1 is requesting the bus, it blocks propagation of grant-signal to other devices.
- Current bus-master indicates to all devices that it is using bus by activating Bus-Busy line (BBSY).
- Arbiter circuit ensures that only one request is granted at any given time according to a predefined priority scheme




Figure 4.20 A simple arrangement for bus arbitration using a daisy chain.

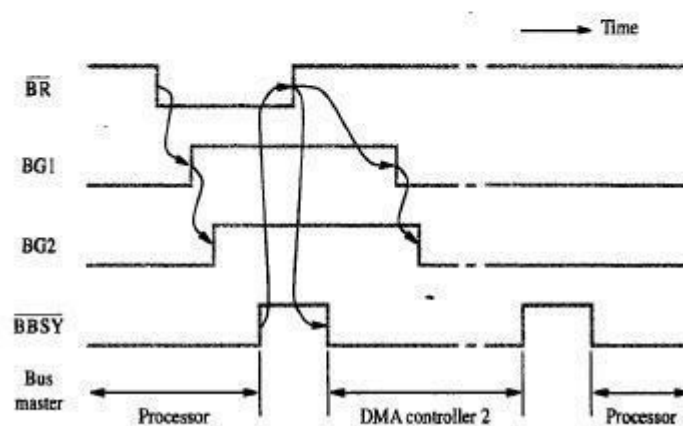


Figure 4.21 Sequence of signals during transfer of bus mastership for the devices in Figure 4.20.

A conflict may arise if both the processor and a DMA controller try to use the bus at the same time to access the main memory. To resolve these conflicts, a special circuit called the bus arbiter is provided to coordinate the activities of all devices requesting memory transfers

## Computer Organization

### DISTRIBUTED ARBITRATION

- All devices participate in the selection of next bus-master (Figure 4.22)
- Each device on bus is assigned a 4-bit identification number (ID).
- When 1 or more devices request bus, they
  - assert Start-Arbitration signal &
  - place their 4-bit ID numbers on four open-collector lines *ARB 0* through *ARB 3*.
- A winner is selected as a result of interaction among signals transmitted over these lines by all contenders.
- Net outcome is that the code on 4 lines represents request that has the highest ID number.
- Main advantage: This approach offers higher reliability since operation of bus is not dependent on any single device.

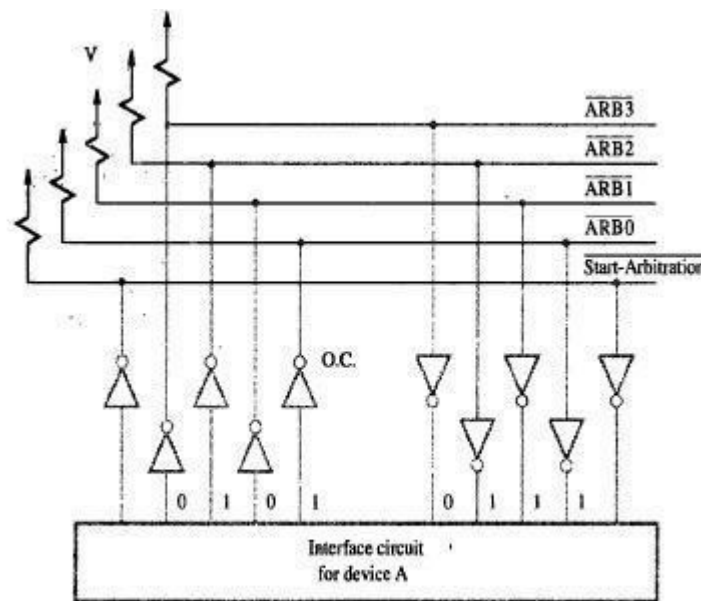


Figure 4.22 A distributed arbitration scheme.